

1. Prove that a symmetric matrix $K \in \mathbb{R}^{d \times d}$ is positive semi-definite if and only if its eigenvalues are nonnegative.
2. Give an example of function $k: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ that is symmetric ($k(x, x') = k(x', x)$) and positive in the sense that $k(x, x') \geq 0$ for all $x, x' \in \mathcal{X}$, but is not positive semidefinite. Conversely, give an example of a kernel that is positive semidefinite, but does not satisfy $k(x, x') \geq 0$ for all $x, x' \in \mathcal{X}$.
3. Given any function $\psi: \mathcal{X} \rightarrow \mathcal{X}'$, prove that if k' is a psd kernel on \mathcal{X}' , then $k(x, x') = k'(\psi(x), \psi(x'))$ is a psd kernel on \mathcal{X} .
4. Prove that if k_1 and k_2 are two positive semi-definite (psd) kernels on a space \mathcal{X} , then
 - (a) the function $k(x, x') := k_1(x, x') + k_2(x, x')$ is a psd kernel on \mathcal{X} ;
 - (b) The function $k_{\oplus}((x_1, x_2), (x'_1, x'_2)) = k_1(x_1, x'_1) + k_2(x_2, x'_2)$ is a psd kernel on $\mathcal{X} \times \mathcal{X}$.
 - (c) Let $\alpha(\mathbf{x}, \mathbf{x}')$ be the angle between $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^n$. Prove that the cosine kernel $k_{\angle}(\mathbf{x}, \mathbf{x}') = \cos \alpha(\mathbf{x}, \mathbf{x}')$ is a psd kernel on $\mathcal{X} = \mathbb{R}^n$.
5. In machine learning one often encounters problems that are invariant to some symmetry group. For example, in learning to recognize objects, if $h(x)$ is the predicted label of object x and $T_{\alpha}(x)$ is the same object rotated by α degrees, one would like h to satisfy $h(T_{\alpha}(x)) = h(x)$ for any α . Formally, this situation corresponds to considering a collection $\mathcal{T}_G = \{T_g \mid g \in G\}$ of transformations $T_g: \mathcal{X} \rightarrow \mathcal{X}$ parametrized by the symmetry group G . For simplicity, let us assume that G is finite.

Given a kernel $k: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$, the “best alignment” similarity measure

$$k_{max}(x, x') = \max_{g, g' \in G} k(T_g(x), T_{g'}(x'))$$

is certainly invariant to \mathcal{T}_G . Show by providing a counterexample that k_{max} is not necessarily positive semi-definite, though. In contrast, show that

$$k_G(x, x') = \sum_{g, g' \in G} k(T_g(x), T_{g'}(x'))$$

is both invariant and positive semi-definite. Show that the output of any kernel based learning algorithm defined in terms of this kernel will satisfy $h(T_g(x)) = h(x)$ for any $g \in G$. Describe the RKHS and feature map corresponding to k_G . If you are into abstract algebra, further show that provided that \mathcal{T}_G is itself a group (as it usually is), then

$$k'_G(x, x') = \sum_{g, g' \in G} k(T_g(x), x')$$

is also an invariant semi-definite kernel and has essentially the same RKHS as k_G .

6. Given a finite alphabet Σ (e.g., $\Sigma = \{a, b, \dots, z\}$), $\bar{x} \in \Sigma^{\ell}$ denotes that $\bar{x} = x_0 x_1 \dots x_{\ell-1}$ is a string of length $\ell_{\bar{x}}$, in which each character x_i comes from Σ . For any increasing sequence $0 \leq i_1 \leq \dots \leq i_k \leq \ell_{\bar{x}} - 1$, we will use $\bar{x}_{(i_1, \dots, i_k)}$ to denote the string $x_{i_1} x_{i_2} \dots x_{i_k}$, and we will use $\bar{x}_{i:j}$ as a shorthand for $x_i x_{i+1} \dots x_j$. Given two characters u and v , we will also use the notation

$$\llbracket u = v \rrbracket = \begin{cases} 1 & \text{if } u = v \\ 0 & \text{if } u \neq v. \end{cases}$$

- (a) Let $n_{\bar{s}}(\bar{x})$ be the number of places that \bar{s} appears in \bar{x} as a contiguous substring, i.e., $n_{\bar{s}}(\bar{x}) = |\{i \mid 0 \leq i \leq \ell_{\bar{x}} - \ell_{\bar{s}}, \bar{x}_{i:i+\ell_{\bar{s}}-1} = \bar{s}\}|$. Define the κ 'th order contiguous substring kernel between two strings \bar{a} and \bar{b} as

$$k_{\kappa}(\bar{a}, \bar{b}) = \sum_{\bar{s} \in \Sigma^{\kappa}} n_{\bar{s}}(\bar{a}) n_{\bar{s}}(\bar{b}).$$

Explain the intuitive meaning of this kernel and why it is positive semi-definite. Write down the pseudocode of an $O(\kappa \ell_{\bar{a}} \ell_{\bar{b}})$ running time algorithm to compute $k_{\kappa}(\bar{a}, \bar{b})$ (here, one unit of computation is testing a single character z against a specific character a_i of \bar{a} or b_j of \bar{b}).

- (b) Let $n_{\bar{s}}^*(\bar{x})$ be the number of places that \bar{s} appears in \bar{x} as a gappy substring, i.e.,

$$n_{\bar{s}}^*(\bar{x}) = |\{(i_1, \dots, i_{\ell_{\bar{s}}}) \mid 0 \leq i_1 < i_2 < \dots < i_{\ell_{\bar{s}}} \leq \ell_{\bar{x}} - 1, \bar{x}_{(i_1, i_2, \dots, i_{\ell_{\bar{s}}})} = \bar{s}\}|.$$

Define the κ 'th order gappy substring kernel between two strings \bar{a} and \bar{b} as

$$k_{\kappa}^*(\bar{a}, \bar{b}) = \sum_{\bar{s} \in \Sigma^{\kappa}} n_{\bar{s}}^*(\bar{a}) n_{\bar{s}}^*(\bar{b}). \quad (1)$$

Explain the intuitive meaning of this kernel and why it is positive semi-definite.

- i. Devise an $O(|\Sigma| + \ell_{\bar{a}} + \ell_{\bar{b}})$ algorithm for computing $k_1^*(\bar{a}, \bar{b})$ and write down its pseudocode.
 - ii. Let $1 \leq p \leq \kappa$. Devise an $O(p)$ algorithm for computing $k_p^*(\bar{a}_{0:p-1}, \bar{b}_{0:p-1})$.
 - iii. Assume that $k_{p-1}^*(\bar{a}_{0:i}, \bar{b}_{0:j})$ and $k_p^*(\bar{a}_{0:i}, \bar{b}_{0:j})$ have been computed for all $0 \leq i \leq N$ and $0 \leq j \leq M$ for some (N, M) .
 - A. For any fixed $j \in \{0, \dots, M\}$, derive a formula for $k_p^*(\bar{a}_{0:N+1}, \bar{b}_{0:j})$ involving only already computed quantities plus the indicators $\llbracket a_{N+1} = b_{p-1} \rrbracket, \dots, \llbracket a_{N+1} = b_j \rrbracket$. Please don't just state the final formula, but also justify it.
 - B. For any fixed $i \in \{0, \dots, N\}$, derive a formula for $k_p^*(\bar{a}_{0:i}, \bar{b}_{0:M+1})$ involving only already computed quantities plus the indicators $\llbracket a_{p-1} = b_{M+1} \rrbracket, \dots, \llbracket a_i = b_{M+1} \rrbracket$. Please don't just state the final formula, but also justify it.
 - C. Combining parts (i)–(iii), derive an $O(\kappa \ell_{\bar{a}} \ell_{\bar{b}})$ algorithm to compute $k_{\kappa}^*(\bar{a}, \bar{b})$ and write down its pseudocode.
- (c) A potential drawback of the gappy substring kernel is that the contribution of any gappy match is counted the same, regardless of whether \bar{s} occurs as a contiguous substring of \bar{a} and \bar{b} or whether it has huge gaps. One solution that has been proposed to this problem is to penalize the match by the number of skipped characters by defining

$$w_{\bar{s}}(\bar{x}) = \sum_{0 \leq i_1 < i_2 < \dots < i_{\ell_{\bar{s}}} \leq \ell_{\bar{x}} - 1} \lambda^{i_{\ell_{\bar{s}}} - i_1 - (\ell_{\bar{s}} - 1)} \llbracket \bar{x}_{(i_1, i_2, \dots, i_{\ell_{\bar{s}}})} = \bar{s} \rrbracket$$

for some penalty factor $\lambda < 1$, and defining the kernel as

$$k_{\kappa}^{\square}(\bar{a}, \bar{b}) = \sum_{\bar{s} \in \Sigma^{\kappa}} w_{\bar{s}}(\bar{a}) w_{\bar{s}}(\bar{b}).$$

Repeat part (b) of this question for k_{κ}^{\square} (part (iii) will have to be somewhat modified) and show that k_{κ}^{\square} can also be computed in time $O(\kappa \ell_{\bar{a}} \ell_{\bar{b}})$.

- (d) **(extra credit)** Another useful variation on the gappy substring kernel is to define a kernel k_{Σ} on the alphabet itself, and, instead of counting exact (gappy) matches between \bar{a} and \bar{b} , define the kernel as

$$k_{\kappa}^{\triangle}(\bar{a}, \bar{b}) = \sum_{0 \leq i_1 < i_2 < \dots < i_{\kappa} \leq \ell_{\bar{a}} - 1} \sum_{0 \leq j_1 < j_2 < \dots < j_{\kappa} \leq \ell_{\bar{b}} - 1} \prod_{t=1}^{\kappa} k_{\Sigma}(a_{i_t}, b_{j_t}).$$

Describe the intuitive meaning of this kernel, why it is positive semi-definite, and show that this kernel can also be computed in time $O(\kappa \ell_{\bar{a}} \ell_{\bar{b}})$.